

# High Performance Computing for Mathematics

## Assignment 1

Jamie Burke

January 30, 2021

### 1 Introduction

The task was to parallelise a matrix multiplication task where the multiplication task was to be split and completed by several different processes using the message passing interface system on Cirrus. Using a procedural language such as FORTRAN, the code was written in order to split up the multiplication task evenly so that each process had to play their part equally in deducing the answer of a simple matrix multiplication problem.

The problem itself was to multiply two matrices,  $A$  and  $B$  producing matrix  $D$ , each of size  $N \times N$ . The matrices are built using the simple rules defined as

$$A_{ij} = (N - j + i + 1)i \quad (1)$$

$$B_{ij} = (j + i)(N - j + 1). \quad (2)$$

As simple as this task is, in order to demonstrate the MP interface the task was to employ a certain number of processors  $P$  such that:

1. The root process,  $P = 0$ , would construct the matrix  $B$  in its entirety. This matrix would be **broadcasted** to all other processes to use for step 3.
2. All other processes (including the root process) would compute a certain number of rows of  $A$  — depending on the size of  $N$  relative to  $P$  — and multiply these rows with matrix  $B$  to deduce their respective rows of the answer  $D$ .
3. The root process would then **gather** all these rows of  $D$  and populate them into the matrix itself to produce and print the final answer.

The caveats here are the sizes of  $N$  and  $P$ . The problem is simple when  $P \geq N$  as we can allocate the first  $N$  processes to computing their individual row of  $A$  and deducing their corresponding row of  $D$ . When  $P < N$  care needs to be taken on how to allocate the processes their rows of  $A$ . If  $N$  is divisible by  $P$ , i.e. for some whole number  $n$  we have  $nP = N$ , then we can allocate  $n$  rows of  $A$  to process  $i$ ,  $i = 0, \dots, P - 1$ , and compute  $n$  rows of  $D$  using matrix  $B$ . This is because we will accumulate the  $N$  rows of  $D$  after all processes have run the code. If  $N$  is not divisible by  $P$  then we compute  $n$  as  $n = \lceil \frac{N}{P} \rceil$  and follow the same procedure for if  $N$  was divisible by  $P$ . However, we will have some junk rows computed that don't form part of the answer in  $D$ . These rows are discarded after gathering all results outputting the correct answer, as seen in the results section.

### 2 Results

This section will show screenshots of tests for various values of  $N$  and  $P$  for the cases where  $N = P$ ,  $N > P$  and  $N < P$ . In particular, choices of  $(N, P) = (3, 3), (4, 6)$  and  $(5, 3)$  will be demonstrated below. In order to validate results, the inbuilt function `matmul()` is used for comparison.

The submission folder contains 3 files, `mat_mul_ext.f90`, `compile_now`, and `mat_mul_job.slurm`. The `.f90` file contains the FORTRAN code detailed in the appendix, `compile_now` contains the bash code to compile the FORTRAN code. The `.slurm` file contains the information to send to Cirrus in order to run the executable outputted from `compile_now`. For further experimentation, the two parameter to change is  $N$  in `mat_mul_ext.f90` and the `tasks-per-node` parameter (equivalent to  $P$  in this report). Below show snapshots of the output file from running the three experiments.

```

Matrix A:
   4      3      2
  10      8      6
  18     15     12

Matrix B:
   6      6      4
   9      8      5
  12     10      6

Multiplying matrices A and B using matmul():
  75      68      43
 204     184     116
 387     348     219

Multiplying matrices A and B using MPI Funcs:
  75      68      43
 204     184     116
 387     348     219

Summed difference between answers:
   0      0      0
   0      0      0
   0      0      0
|

```

Figure 1: Correct results when  $(N, P) = (3, 3)$ .

```

Matrix A:
   5      4      3      2
  12     10      8      6
  21     18     15     12
  32     28     24     20

Matrix B:
   8      9      8      5
  12     12     10      6
  16     15     12      7
  20     18     14      8

Multiplying matrices A and B using matmul():
 176     174     144      86
 464     456     376     224
 864     846     696     414
1376     1344     1104     656

Multiplying matrices A and B using MPI Funcs:
 176     174     144      86
 464     456     376     224
 864     846     696     414
1376     1344     1104     656

Summed difference between answers:
   0      0      0      0
   0      0      0      0
   0      0      0      0
   0      0      0      0
|

```

Figure 2: Correct results when  $(N, P) = (4, 6)$ .

```

Matrix A:
  6      5      4      3      2
 14     12     10     8      6
 24     21     18     15     12
 36     32     28     24     20
 50     45     40     35     30

Matrix B:
 10     12     12     10     6
 15     16     15     12     7
 20     20     18     14     8
 25     24     21     16     9
 30     28     24     18     10

Multiplying matrices A and B using matmul():
 350     360     330     260     150
 900     920     840     660     380
1650    1680    1530    1200     690
2600    2640    2400    1880    1080
3750    3800    3450    2700    1550

Multiplying matrices A and B using MPI Funcs:
 350     360     330     260     150
 900     920     840     660     380
1650    1680    1530    1200     690
2600    2640    2400    1880    1080
3750    3800    3450    2700    1550

Summed difference between answers:
  0      0      0      0      0
  0      0      0      0      0
  0      0      0      0      0
  0      0      0      0      0
  0      0      0      0      0

```

Figure 3: Correct results when  $(N, P) = (5, 3)$ .

## A FORTRAN Code

```
1 !-----
2 ! Perform matrix multiplication using MPI functions.
3 ! This program is able to deal with any combination of
4 ! number of processes and size of matrices being multiplied
5 !-----
6     program mat_mul_ext
7         implicit none
8
9         ! Include the mpif.h module for MPI function
10        include "mpif.h"
11
12        ! Define size of matrices and iterators i, j
13        integer :: N, i, j
14        parameter(N = 3)
15        ! Define the integer variables necessary for MPI functions
16        integer :: comm, rank, N_processes, ierr
17
18        ! Define matrices A, B, C_builtin, D and vectors
19        ! row_c and row_a (to hold interim results of MPI computation
20        ! note: C_builtin will hold the computation using matmul()
21        ! while D will hold the result computed using MPI functions
22        integer, dimension(N, N) :: A, B, C_builtin, CD_diff, D
23        integer, dimension(N) :: row_a, row_c
24
25        ! Create allocatable 2D arrays for when N_processes < N
26        ! Also create integer status variable for allocation
27        integer :: errCode
28        integer, dimension(:, :), allocatable :: rows_a, rows_c
29        integer, dimension(:, :), allocatable :: interim_D
30
31        ! define remainder between the size of matrices and
32        ! the number of processors. Also define divisor, the
33        ! whole integer number when dividing N by N_processes
34        integer :: div, rem
35
36        ! Initialise MPI and extract the rank and size of process
37        comm = MPI_COMM_WORLD
38        call MPI_INIT(ierr)
39        call MPI_COMM_RANK(comm, rank, ierr)
40        call MPI_COMM_SIZE(comm, N_processes, ierr)
41
42        ! Set up conditions depending on rank (process)
43        if (rank == 0) then
44            ! Create matrix B if rank 0 and form 1st row of A
45            do i = 1, N
46                do j = 1, N
47                    B(i, j) = (j + i) * (N - j + 1)
48                end do
49            end do
50        end if
51
52        ! Broadcast B to all other processes
53        call MPI_BCAST(B, N**2, MPI_INT, 0, &
54                     comm, ierr)
55
56        ! Compute divisor between N and N_processes for use
57        ! when number of processes is less than the size of
```

```

58      ! matrices
59      rem = mod(N, N_processes)
60      div = floor(float(N)/float(N_processes))
61
62      ! If the number of processes is greater than or equal
63      ! to the size of matrices then only use up to N of the
64      ! processes in dealing with the computation
65      if (N_processes >= N) then
66          if (rank < N) then
67              ! Form rank'th row of A
68              do j = 1, N
69                  row_a(j) = (N - j + (rank+1) + 1)*(rank + 1)
70              end do
71
72              ! Compute vector c, i.e. compute matmul(row_a,B)
73              row_c = matmul(row_a, B)
74          end if
75
76          ! Gather vectors rows_c to put in matrix D
77          call MPI_GATHER(row_c, N, MPI_INT, D, N, &
78                        MPI_INT, 0, comm, ierr)
79
80      ! If the number of the processes are less than N, assign an
81      ! equal number of rows of A to each process and populate D.
82      else if (N_processes < N) then
83          ! If non-zero remainder (N is not divisible by N_processes
84          ! then add 1 to the divisor so compute ceil(N/N_processes)
85          if (rem /= 0) then
86              div = div + 1
87          end if
88
89          ! Allocate the necessary memory to rows_a and rows_c
90          ! Also allocate maximum number of columns to interim_D
91          allocate(rows_a(div, N), stat = errCode)
92          allocate(rows_c(div, N), stat = errCode)
93          allocate(interim_D(N, div*N_processes), stat = errCode)
94
95          ! Build rows of A. The row number is maintained by
96          ! iterating over the divisor between N and N_processes and
97          ! the rank of the process.
98          do i = 1, div
99              do j = 1, N
100                  rows_a(i, j) = (N - j + (div*rank+i) + 1)*(div*rank+i)
101              end do
102          end do
103
104          ! Compute rows of answer and transpose before gathering
105          ! them to populate rows of D
106          rows_c = matmul(rows_a, B)
107          rows_c = transpose(rows_c)
108
109          ! Gather vector rows_c to populate D. If N is not divisible
110          ! by number of processes then populate interim_D and
111          ! discard the last few columns (since FORTRAN is a
112          ! column-wise language) which correspond to the extra rows
113          ! computed.
114          if (rem /= 0) then
115              call MPI_GATHER(rows_c, div*N, MPI_INT, interim_D, &
116                            div*N, MPI_INT, 0, comm, ierr)

```

```

117         D = interim_D(:, :N)
118     else
119         call MPI_GATHER(rows_c, div*N, MPI_INT, D, div*N, &
120             MPI_INT, 0, comm, ierr)
121     end if
122 end if
123
124 ! Print results
125 if (rank == 0) then
126
127     ! Loop over iterators to define matrices A and B
128     do i = 1, N
129         do j = 1, N
130             A(i, j) = (N - j + i + 1) * i
131         end do
132     end do
133
134     print*, "Matrix A:"
135     do i = 1, N
136         print*, A(i, :)
137     end do
138
139     print*, "Matrix B:"
140     do i = 1, N
141         print*, B(i, :)
142     end do
143
144     print*, " "
145     C_builtin = matmul(A, B)
146     print*, "Multiplying matrices A and B using matmul():"
147     do i = 1, N
148         print*, C_builtin(i, :)
149     end do
150
151     print*, " "
152     print*, "Multiplying matrices A and B using MPI Funcs:"
153     D = transpose(D)
154     do i = 1, N
155         print*, D(i, :)
156     end do
157
158     print*, " "
159     CD_diff = C_builtin - D
160     print*, "Summed difference between answers:"
161     do i = 1, N
162         print*, CD_diff(i, :)
163     end do
164
165 end if
166
167 ! Stop MPI process
168 call MPI_FINALIZE(ierr)
169
170 end program mat_mul_ext

```